

# Debugging Techniques for Q-Learning in C++

## Introduction

If after making adjustments, the agent's behavior still does not change, the problem might be in the following areas:

- **Q-Table not updating:** The agent might be stuck in some loop state or strategy, preventing the Q-Table from effectively updating.
- **Insufficient exploration:** If the agent has adopted a "lazy" strategy and the epsilon value is too low, it may always choose the same action, preventing it from exploring better paths.
- **Action not updated:** Although the code updates the Q-Table and selects actions, the actual action selection may not be reflected in the agent's behavior.

## Deep Debugging Suggestions

### 1. Print Q-Table

Print the contents of the Q-Table at the end of each episode to see if the Q-values for state-action pairs are being updated, especially those that cause the agent to stay in one place.

### 2. Increase Exploration Rate

Increase the epsilon value to raise the probability of exploration and see if the agent tries different actions. You can add a conditional check in the `choose_action()` function to force exploration.

### 3. Output Debugging Information

Add more debug information in the `step()` function and the Q-Table update logic to check the state, action, reward, and Q-value changes at each step.

## Detailed Debugging Code

### Modify choose\_action Function

```
int choose_action(int state) {
    if ((double)rand() / RAND_MAX < EPSILON) { // If random number is less than
        int random_action = rand() % NUM_ACTIONS; // Randomly choose an action
        std::cout << "Exploring:_" << random_action << "_for_state_" << state << "\n";
        return random_action;
    } else { // Otherwise exploit
        int best_action = std::max_element(Q_table[state].begin(), Q_table[state].end(),
            // Choose action with max Q-value
            [&state, &Q_table](int a, int b) { return Q_table[state][a] > Q_table[state][b]; })->first;
        std::cout << "Exploiting:_" << best_action << "_for_state_" << state << "\n";
        return best_action;
    }
}
```

### Modify update\_Q Function

```
void update_Q(int state, int action, int reward, int next_state) {
    double best_next_action = *std::max_element(Q_table[next_state].begin(), Q_table[next_state].end(),
        // Find max Q-value for next state
        [&Q_table](int a, int b) { return Q_table[next_state][a] > Q_table[next_state][b]; });
    double old_value = Q_table[state][action];
    Q_table[state][action] += ALPHA * (reward + GAMMA * best_next_action - Q_table[state][action]);
    // Update Q-value using Q-Learning formula
    std::cout << "Updated_Q(" << state << ", " << action << ")_from_" << old_value << " to " << Q_table[state][action] << "\n";
}
```

### Print Q-Table Contents

```
void print_Q_table() {
    std::cout << "Q-table:" << std::endl;
    for (size_t i = 0; i < Q_table.size(); ++i) {
        std::cout << "State_" << i << ":\n";
        for (size_t j = 0; j < Q_table[i].size(); ++j) {
            std::cout << Q_table[i][j] << " ";
        }
        std::cout << std::endl;
    }
}
```

### Call Debug Functions in Main Loop

```

for (int episode = 0; episode < 1000; ++episode) {
    env.reset(); // Reset environment
    int state = env.get_state(); // Get current state
    int total_reward = 0; // Initialize total reward to 0

    while (!env.is_done()) { // While task is not done
        int action = agent.choose_action(state); // Agent chooses action
        int reward = env.step(action); // Execute action, get reward
        int next_state = env.get_state(); // Get next state
        agent.update_Q(state, action, reward, next_state);
    // Update Q-table

        state = next_state; // Update current state to next state
        total_reward += reward; // Accumulate reward for this episode
    }

    std::cout << "Episode_" << episode << ",_Total_Reward:_" << total_reward <<
    env.render(); // Render environment to show agent and goal

    if (episode % 10 == 0) {
        agent.print_Q_table(); // Print Q-table every 10 episodes
    }
}

```

## Analyze Debug Output

- Check if the Q-Table is being updated during training, particularly for states leading the agent to stay in one place.
- Observe the action selection output to see if there is too much "exploitation," causing the agent to always choose the same action.
- Through these debug messages, identify and fix the issue preventing the agent from moving.

## Other Considerations

- If some Q-values in the Q-Table are too high, they may cause the agent to always choose those actions. Consider introducing a mechanism to periodically reset or smooth the Q-Table to avoid overly confident decisions.
- If the above methods are still ineffective, try step-by-step debugging to ensure each part is functioning as expected.